

Week 3A

Pricilla

2025-09-05

I used a CSV that I downloaded from <https://www.kaggle.com/datasets/sunnykusawa/stock-price-dataset-eod> to create a stock_price_data table in PGADMIN. This data is from 2020. It is 1600 rows of stock prices over a duration of 5 months.

```
library(RSQLite)
library(DBI)
library(sqldf)
```

```
## Loading required package: gsubfn
```

```
## Loading required package: proto
```

```
library(RODBC)
library(odbc)
library(crayon)
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(RPostgres)
library(lubridate)
```

```
##
```

```
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      date, intersect, setdiff, union
```

```
library(ggplot2)
```

```
##  
## Attaching package: 'ggplot2'  
  
## The following object is masked from 'package:crayon':  
##  
##      %+%
```

1. Find a dataset that includes time series for two or more separate items. For example, you could use end of day stock or cryptocurrency prices since Jan 1, 2022 for several instruments.

Connected to PgAmin to create Stock_df.

```
# Connect to the database  
con <- dbConnect(odbc::odbc(), "Post")  
  
# Run your query  
Stock <- dbGetQuery(con, "SELECT * FROM stock_price_data;")
```

I will be using the adj_close for my window functions.

```
stock_df <- Stock  
glimpse(stock_df)
```

```
## Rows: 1,600  
## Columns: 13  
## $ open      <dbl> 1737.27, 1765.00, 1787.23, 1744.91, 1729.00, 1723.93, 1729.~  
## $ high      <dbl> 1757.50, 1767.76, 1788.47, 1787.00, 1742.41, 1744.11, 1732.~  
## $ low       <dbl> 1736.09, 1728.00, 1755.11, 1741.82, 1724.35, 1721.20, 1705.~  
## $ close     <dbl> 1752.64, 1736.25, 1757.76, 1773.96, 1734.16, 1728.23, 1720.~  
## $ volume    <dbl> 1053479, 1051308, 986287, 1379642, 465638, 1148684, 1018829~  
## $ adj_high  <dbl> 1757.50, 1767.76, 1788.47, 1787.00, 1742.41, 1744.11, 1732.~  
## $ adj_low   <dbl> 1736.09, 1728.00, 1755.11, 1741.82, 1724.35, 1721.20, 1705.~  
## $ adj_close <dbl> 1752.64, 1736.25, 1757.76, 1773.96, 1734.16, 1728.23, 1720.~  
## $ adj_open  <dbl> 1737.27, 1765.00, 1787.23, 1744.91, 1729.00, 1723.93, 1729.~  
## $ adj_volume <dbl> 1053479, 1051308, 986287, 1379642, 465638, 1148684, 1018829~  
## $ symbol    <chr> "GOOGL", "GOOGL", "GOOGL", "GOOGL", "GOOGL", "GOOGL", "GOOG~  
## $ exchange  <chr> "XNAS", "XNAS", "XNAS", "XNAS", "XNAS", "XNAS", "XNAS", "XN~  
## $ date      <date> 2020-12-31, 2020-12-30, 2020-12-29, 2020-12-28, 2020-12-24~
```


Used SQL to find Average close cost by the year to see which stock did well.

```
YearlyAvg <- dbGetQuery(con, "  
  SELECT  
    symbol,  
    EXTRACT(YEAR FROM date) AS year,  
    AVG(adj_close) AS avg_adj_close  
  FROM stock_price_data  
  GROUP BY symbol, EXTRACT(YEAR FROM date)  
  ORDER BY symbol, year  
")
```

```
glimpse(YearlyAvg)
```

```
## Rows: 16
## Columns: 3
## $ symbol      <chr> "AAPL", "ABT", "AMZN", "BABA", "DIS", "FB", "GOOGL", "JN~
## $ year        <dbl> 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020, 20~
## $ avg_adj_close <dbl> 119.0005, 107.1233, 3198.7356, 276.2280, 138.6336, 272.3~
```

```
# Plot
ggplot(YearlyAvg, aes(x = reorder(symbol, avg_adj_close), y = avg_adj_close, fill = symbol)) +
  geom_col() +
  coord_flip() + # Flip for easier readability
  labs(
    title = "Average Adjusted Close Price by Stock (2020)",
    x = "Stock Symbol",
    y = "Average Adjusted Close Price"
  ) +
  theme_minimal() +
  theme(legend.position = "none")
```



AMAZON AND GOOGLE are the best performing stocks in this dataset.

2. Use window functions (in SQL) to calculate the year-to-date average

```
WinYrAvg <- dbGetQuery(con, "
```

```
SELECT
  symbol,
  date,
  EXTRACT(YEAR FROM date) AS year,
  adj_close,
  AVG(adj_close) OVER (
    PARTITION BY symbol, EXTRACT(YEAR FROM date)
    ORDER BY date
    ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
  ) AS ytd_avg_adj_close
FROM stock_price_data
")
```

```
glimpse(WinYrAvg)
```

```
## Rows: 1,600
## Columns: 5
## $ symbol      <chr> "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAP~
## $ date        <date> 2020-08-11, 2020-08-12, 2020-08-13, 2020-08-14, 202~
## $ year        <dbl> 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020~
## $ adj_close    <dbl> 109.19, 112.82, 114.81, 114.71, 114.41, 115.36, 115.~
## $ ytd_avg_adj_close <dbl> 109.1900, 111.0050, 112.2733, 112.8825, 113.1880, 11~
```

```
WinYrAvg %>%
  filter(symbol %in% c("GOOGL", "AMZN")) %>%
  group_by(symbol) %>%
  slice_head(n = 10)
```

```
## # A tibble: 20 x 5
## # Groups:   symbol [2]
##   symbol date      year adj_close ytd_avg_adj_close
##   <chr> <date>    <dbl>    <dbl>    <dbl>
## 1 AMZN  2020-08-11  2020    3081.    3081.
## 2 AMZN  2020-08-12  2020    3162.    3121.
## 3 AMZN  2020-08-13  2020    3161.    3135.
## 4 AMZN  2020-08-14  2020    3148.    3138.
## 5 AMZN  2020-08-17  2020    3182.    3147.
## 6 AMZN  2020-08-18  2020    3312.    3174.
## 7 AMZN  2020-08-19  2020    3260.    3187.
## 8 AMZN  2020-08-20  2020    3297.    3201.
## 9 AMZN  2020-08-21  2020    3285.    3210.
## 10 AMZN 2020-08-24  2020    3307.    3220.
## 11 GOOGL 2020-08-11  2020    1481.    1481.
## 12 GOOGL 2020-08-12  2020    1507.    1494.
## 13 GOOGL 2020-08-13  2020    1517.    1501.
## 14 GOOGL 2020-08-14  2020    1505.    1502.
## 15 GOOGL 2020-08-17  2020    1516.    1505.
## 16 GOOGL 2020-08-18  2020    1556.    1514.
## 17 GOOGL 2020-08-19  2020    1545.    1518.
## 18 GOOGL 2020-08-20  2020    1576.    1525.
```

```
## 19 GOOGL 2020-08-21 2020 1576. 1531.
## 20 GOOGL 2020-08-24 2020 1585. 1536.
```

The six-day moving averages for each item..

```
Day6Avg <- dbGetQuery(con, "
SELECT
  symbol,
  date,
  adj_close,
  AVG(adj_close) OVER (
    PARTITION BY symbol
    ORDER BY date
    ROWS BETWEEN 5 PRECEDING AND CURRENT ROW
  ) AS day6_adj_close
FROM stock_price_data
ORDER BY symbol, date
")
```

```
glimpse(Day6Avg)
```

```
## Rows: 1,600
## Columns: 4
## $ symbol      <chr> "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", ~
## $ date        <date> 2020-08-11, 2020-08-12, 2020-08-13, 2020-08-14, 2020-0~
## $ adj_close   <dbl> 109.19, 112.82, 114.81, 114.71, 114.41, 115.36, 115.51, ~
## $ day6_adj_close <dbl> 109.1900, 111.0050, 112.2733, 112.8825, 113.1880, 113.5~
```

```
Day6Avg %>%
  filter(symbol %in% c("GOOGL", "AMZN")) %>%
  group_by(symbol) %>%
  slice_head(n = 10)
```

```
## # A tibble: 20 x 4
## # Groups:   symbol [2]
##   symbol date      adj_close day6_adj_close
##   <chr> <date>      <dbl>      <dbl>
## 1 AMZN 2020-08-11    3081.      3081.
## 2 AMZN 2020-08-12    3162.      3121.
## 3 AMZN 2020-08-13    3161.      3135.
## 4 AMZN 2020-08-14    3148.      3138.
## 5 AMZN 2020-08-17    3182.      3147.
## 6 AMZN 2020-08-18    3312.      3174.
## 7 AMZN 2020-08-19    3260.      3204.
## 8 AMZN 2020-08-20    3297.      3227.
## 9 AMZN 2020-08-21    3285.      3248.
## 10 AMZN 2020-08-24    3307.      3274.
## 11 GOOGL 2020-08-11    1481.      1481.
## 12 GOOGL 2020-08-12    1507.      1494.
## 13 GOOGL 2020-08-13    1517.      1501.
## 14 GOOGL 2020-08-14    1505.      1502.
```

```
## 15 GOOGL 2020-08-17 1516. 1505.
## 16 GOOGL 2020-08-18 1556. 1514.
## 17 GOOGL 2020-08-19 1545. 1524.
## 18 GOOGL 2020-08-20 1576. 1536.
## 19 GOOGL 2020-08-21 1576. 1546.
## 20 GOOGL 2020-08-24 1585. 1559.
```


SNAPSHOT OF 6 DAY MOVING AVG FOR AAPL, AMZN & GOOGL.

The 6 day Moving Average Smooths volatility and reduces “noise” from daily price swings.

It helps highlight the general direction (upward/downward) over time.

Traders often use moving averages to identify buy/sell points (e.g., when price crosses the average line).

```
Day6Avg %>%
  filter(symbol %in% c("AAPL", "AMZN", "GOOGL")) %>% # Filter for some symbols to keep plot readable
  ggplot(aes(x = date)) +
  geom_line(aes(y = adj_close, color = "Adjusted Close")) +
  geom_line(aes(y = day6_adj_close, color = "6-Day Avg Adj Close")) +
  facet_wrap(~symbol, scales = "free_y") +
  labs(
    title = "Adjusted Close vs 6-Day Average Adjusted Close",
    y = "Price (USD)",
    color = "Legend"
  ) +
  theme_minimal()
```



AAPL

The blue and red lines follow each other closely, but you can see when the blue line spikes or dips suddenly, the red line moves more gradually.

AMZN & GOOGL are Similar, smoothing is visible, the red line delays in reacting to sharp blue line movements, making it useful for identifying trends rather than daily changes.

Use window functions (dplyr) to calculate the year-to-date average

```
YearAVGdf <- stock_df %>%
  mutate(year = year(date)) %>%           # Extract year from date
  arrange(symbol, date) %>%             # Ensure proper order for cumulative average
  group_by(symbol, year) %>%
  mutate(
    ytd_avg_adj_close = cummean(adj_close) # Cumulative average within each symbol-year
  ) %>%
  ungroup()
```

```
YearAVGdiff <- stock_df %>%
  mutate(year = lubridate::year(date)) %>% # Extract year from date (ensure lubridate is loaded)
  arrange(symbol, date) %>%             # Ensure proper order for cumulative average
  group_by(symbol, year) %>%
  mutate(
    ytd_avg_adj_close = cummean(adj_close), # Cumulative average within each symbol-year
    diff_from_avg = adj_close - ytd_avg_adj_close # Difference between actual and average
  ) %>%
  ungroup()
```

```
glimpse(YearAVGdf)
```

```
## Rows: 1,600
## Columns: 15
## $ open      <dbl> 447.88, 441.99, 457.72, 459.32, 464.25, 457.41, 463.~
## $ high      <dbl> 449.93, 453.10, 464.17, 460.00, 464.35, 464.00, 468.~
## $ low       <dbl> 436.43, 441.19, 455.71, 452.18, 455.85, 456.03, 462.~
## $ close     <dbl> 437.50, 452.04, 460.04, 459.63, 458.43, 462.25, 462.~
## $ volume    <dbl> 46975594, 41486205, 52520516, 41391302, 29431414, 26~
## $ adj_high  <dbl> 112.29, 113.08, 115.84, 114.80, 115.89, 115.80, 116.~
## $ adj_low   <dbl> 108.92, 110.11, 113.73, 112.85, 113.77, 113.81, 115.~
## $ adj_close <dbl> 109.19, 112.82, 114.81, 114.71, 114.41, 115.36, 115.~
## $ adj_open  <dbl> 111.78, 110.31, 114.23, 114.63, 115.86, 114.16, 115.~
## $ adj_volume <dbl> 187902376, 165944820, 210082064, 165565208, 11772565~
## $ symbol    <chr> "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAPL", "AAP~
## $ exchange  <chr> "XNAS", "XNAS", "XNAS", "XNAS", "XNAS", "XNAS", "XNA~
## $ date      <date> 2020-08-11, 2020-08-12, 2020-08-13, 2020-08-14, 202~
## $ year      <dbl> 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020, 2020~
## $ ytd_avg_adj_close <dbl> 109.1900, 111.0050, 112.2733, 112.8825, 113.1880, 11~
```

```
YearAVGdf %>%
  filter(symbol %in% c("GOOGL", "AMZN")) %>%
  group_by(symbol) %>%
  slice_head(n = 10)
```

```
## # A tibble: 20 x 15
## # Groups:   symbol [2]
##   open  high   low close volume adj_high adj_low adj_close adj_open adj_volume
##   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl>  <dbl>   <dbl>   <dbl>   <dbl>
## 1 3113. 3159. 3073 3081. 3.72e6 3159. 3073   3081. 3113. 3718057
## 2 3108 3174. 3101. 3162. 3.53e6 3174. 3101. 3162. 3108 3527229
## 3 3183. 3218. 3155 3161. 3.15e6 3218. 3155   3161. 3183. 3149043
## 4 3178. 3178. 3120 3148. 2.75e6 3178. 3120   3148. 3178. 2751723
## 5 3173. 3195. 3154. 3182. 2.68e6 3195. 3154. 3182. 3173. 2675890
## 6 3212 3320 3206. 3312. 5.35e6 3320 3206. 3312. 3212 5345974
## 7 3303. 3316. 3256 3260. 4.19e6 3316. 3256   3260. 3303. 4185137
## 8 3252 3313. 3238 3297. 3.33e6 3313. 3238   3297. 3252 3332478
## 9 3295 3314. 3275. 3285. 3.58e6 3314. 3275. 3285. 3295 3575862
## 10 3310. 3380. 3258. 3307. 4.67e6 3380. 3258. 3307. 3310. 4666258
## 11 1494 1510. 1478. 1481. 1.55e6 1510. 1478. 1481. 1494 1554853
## 12 1487. 1512. 1485 1507. 1.13e6 1512. 1485   1507. 1487. 1126560
## 13 1508. 1537. 1508. 1517. 1.12e6 1537. 1508. 1517. 1508. 1119711
## 14 1514. 1520. 1499 1505. 1.10e6 1520. 1499   1505. 1514. 1097059
## 15 1516. 1524. 1505 1516. 9.98e5 1524. 1505   1516. 1516. 998010
## 16 1526. 1557. 1522. 1556. 1.42e6 1557. 1522. 1556. 1526. 1418850
## 17 1552. 1569. 1540 1545. 1.52e6 1569. 1540   1545. 1552. 1523592
## 18 1540. 1580. 1534. 1576. 1.32e6 1580. 1534. 1576. 1540. 1319131
## 19 1572. 1592. 1562. 1576. 1.74e6 1592. 1562. 1576. 1572. 1742275
## 20 1592. 1609. 1575. 1585. 1.28e6 1609. 1575. 1585. 1592. 1281893
## # i 5 more variables: symbol <chr>, exchange <chr>, date <date>, year <dbl>,
## #   ytd_avg_adj_close <dbl>
```

```
YearAVGdf %>%
  filter(symbol %in% c("GOOGL", "AMZN")) %>%
  group_by(symbol) %>%
  slice_head(n = 10) %>%
  select(symbol, date, adj_close, ytd_avg_adj_close)
```

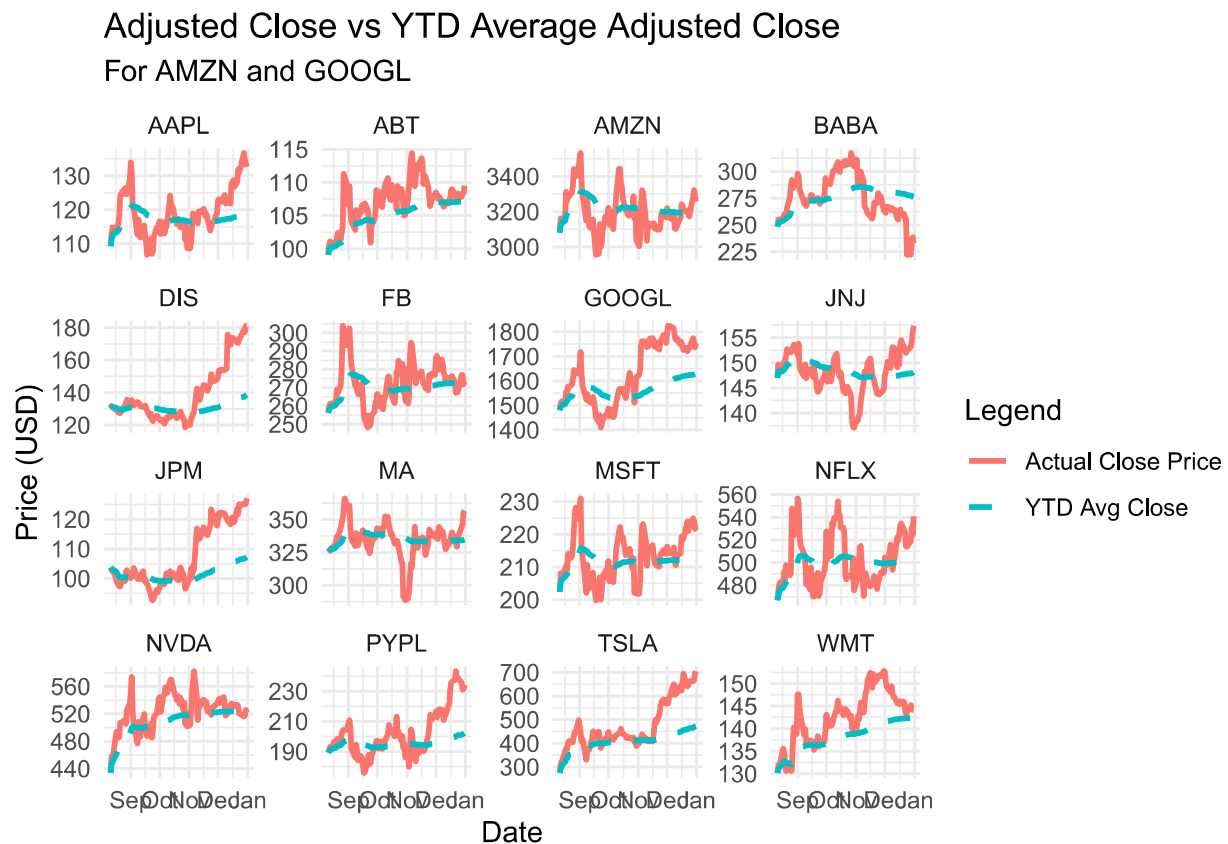
```
## # A tibble: 20 x 4
## # Groups:   symbol [2]
##   symbol date      adj_close ytd_avg_adj_close
##   <chr> <date>      <dbl>      <dbl>
## 1 AMZN 2020-08-11 3081.      3081.
## 2 AMZN 2020-08-12 3162.      3121.
## 3 AMZN 2020-08-13 3161.      3135.
## 4 AMZN 2020-08-14 3148.      3138.
## 5 AMZN 2020-08-17 3182.      3147.
## 6 AMZN 2020-08-18 3312.      3174.
## 7 AMZN 2020-08-19 3260.      3187.
## 8 AMZN 2020-08-20 3297.      3201.
## 9 AMZN 2020-08-21 3285.      3210.
## 10 AMZN 2020-08-24 3307.      3220.
## 11 GOOGL 2020-08-11 1481.      1481.
## 12 GOOGL 2020-08-12 1507.      1494.
## 13 GOOGL 2020-08-13 1517.      1501.
## 14 GOOGL 2020-08-14 1505.      1502.
## 15 GOOGL 2020-08-17 1516.      1505.
## 16 GOOGL 2020-08-18 1556.      1514.
## 17 GOOGL 2020-08-19 1545.      1518.
```



```
## 18 GOOGL 2020-08-20 1576. 1525.
## 19 GOOGL 2020-08-21 1576. 1531.
## 20 GOOGL 2020-08-24 1585. 1536.
```

```
ggplot(YearAVGdf, aes(x = date)) +
  geom_line(aes(y = adj_close, color = "Actual Close Price"), size = 1) +
  geom_line(aes(y = ytd_avg_adj_close, color = "YTD Avg Close"), linetype = "dashed", size = 1) +
  facet_wrap(~symbol, scales = "free_y") + # One panel per stock
  labs(
    title = "Adjusted Close vs YTD Average Adjusted Close",
    subtitle = "For AMZN and GOOGL",
    x = "Date",
    y = "Price (USD)",
    color = "Legend"
  ) +
  theme_minimal()
```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use 'linewidth' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```



Three stock snapshot Explanation

AMZN (Amazon)

Red Line (Actual): Shows significant fluctuations with a general upward trend.

Blue Line (YTD Average): Smoothly rises over time, reflecting Amazon's growing performance throughout the year.

The stock had sharp peaks and dips, likely due to market volatility or earnings releases.

Despite volatility, the YTD average increases steadily, suggesting consistent overall growth.

Amazon consistently traded above the YTD average, indicating strong investor confidence.

Google

Red Line and Actual: Starts lower and gradually increases with fewer sharp fluctuations compared to AMZN.

Blue Line: Steady linear climb, indicating a healthy and consistent rise in adjusted close prices.

GOOGL experienced less volatility than AMZN.

The price often returned to the YTD average, showing alignment between market price and intrinsic valuation.

Late in the year, the actual price diverged above the average signalling a strong bullish signal.

TSLA-Tesla

Red Line Actual price show high volatility with steep increases and dips.

The Blue Line which is YTD is Steady but much lower than the actual price for most of the timeline.

Tesla's stock saw explosive growth, which isn't matched by the YTD average until late.

The gap between actual price and YTD average indicates rapid, recent price surges.

Ideal example of a high momentum stock with elevated risk.

When I Calculate the difference between Average Price and adj_close it makes sense to not time the market in hopes of getting a good deal but to just continue buying stock because for my favorite stock AMZN I am not seeing a significant dip in prices that warrants timing the market for a good deal. 2020-08-18 to 2020-08-24 would have been a good day to sell because the average Price was high.

```
YearAVGdiff %>%
  filter(symbol %in% c("GOOGL", "AMZN")) %>%
  group_by(symbol) %>%
  slice_head(n = 10) %>%
  select(symbol, date, adj_close, ytd_avg_adj_close, diff_from_avg)
```

```
## # A tibble: 20 x 5
## # Groups:   symbol [2]
##   symbol date      adj_close ytd_avg_adj_close diff_from_avg
##   <chr> <date>      <dbl>      <dbl>      <dbl>
## 1 AMZN  2020-08-11    3081.        3081.         0
## 2 AMZN  2020-08-12    3162.        3121.        40.8
## 3 AMZN  2020-08-13    3161.        3135.        26.4
## 4 AMZN  2020-08-14    3148.        3138.        10.0
## 5 AMZN  2020-08-17    3182.        3147.        35.5
## 6 AMZN  2020-08-18    3312.        3174.       138.
## 7 AMZN  2020-08-19    3260.        3187.        73.7
## 8 AMZN  2020-08-20    3297.        3201.        96.8
## 9 AMZN  2020-08-21    3285.        3210.        74.8
## 10 AMZN 2020-08-24    3307.        3220.        87.8
## 11 GOOGL 2020-08-11    1481.        1481.         0
```

```
## 12 GOOGL 2020-08-12 1507. 1494. 13.4
## 13 GOOGL 2020-08-13 1517. 1501. 15.2
## 14 GOOGL 2020-08-14 1505. 1502. 2.37
## 15 GOOGL 2020-08-17 1516. 1505. 11.2
## 16 GOOGL 2020-08-18 1556. 1514. 42.3
## 17 GOOGL 2020-08-19 1545. 1518. 26.7
## 18 GOOGL 2020-08-20 1576. 1525. 51.0
## 19 GOOGL 2020-08-21 1576. 1531. 44.7
## 20 GOOGL 2020-08-24 1585. 1536. 48.9
```

```
# Filter only AMZN and GOOGL
filtered_df <- YearAVGdiff %>%
  filter(symbol %in% c("AMZN", "GOOGL"))


# Plot: Actual vs YTD Avg Close
ggplot(filtered_df, aes(x = date)) +
  geom_line(aes(y = adj_close, color = "Actual Close Price"), size = 1) +
  geom_line(aes(y = ytd_avg_adj_close, color = "YTD Avg Close"), size = 1, linetype = "dashed") +
  facet_wrap(~symbol, scales = "free_y") +
  labs(
    title = "Adjusted Close vs YTD Average Close",
    x = "Date", y = "Price (USD)",
    color = "Legend"
  ) +
  theme_minimal() +
  scale_color_manual(values = c("Actual Close Price" = "red", "YTD Avg Close" = "blue"))
```

Adjusted Close vs YTD Average Close



Plot for how far the actual close was from the YTD average This Plot shows the difference between the daily Adjusted Close price and the Year-to-Date (YTD) Average Adjusted Close for two stocks. AMZN (Amazon) and GOOGL over the latter part of 2020.

```
ggplot(filtered_df, aes(x = date, y = diff_from_avg, fill = symbol)) +  
  geom_col(position = "dodge") +  
  labs(  
    title = "Difference from YTD Average Adjusted Close",  
    x = "Date", y = "Difference (USD)"  
  ) +  
  facet_wrap(~symbol, scales = "free_y") +  
  theme_minimal()
```



AMZN

The differences swing widely between +250 and -300 USD showing high volatility.

In early August, AMZN traded above its YTD average.

In late September to mid-October, AMZN consistently traded below the YTD average — suggesting a slump relative to its average performance.

Later, it bounced back but still fluctuated around the YTD average.

GOOGL

Google showed a more consistent uptrend compared to AMZN.

After some minor dips in August and September, GOOGL showed mostly positive differences. Closing well above its YTD average for most of November and December.

This suggests strong, steady performance in the second half of the year.